

How to work on course projects

Every project is divided into 2 parts. Part 1 of every project is a quiz. The quizzes need to be answered based on data exploration of the dataset given and some generic questions about algorithms discussed in the course. Hence before attempting the quiz you need to conduct exploratory analysis of the data along with data cleaning. For the quizzes, you need to consider only the training dataset. Part 2 is where you create the machine learning model using the exploratory analysis and data cleaning you have done for part 1. More detailed explanation of what needs to be done for each project is available in each of the project descriptions. The below guide is for you to understand how to approach and work on a project and is just an example.

In this guide we will discuss how to work on a sample classification problem. We are not discussing problem details here in terms of what all variables explicitly mean, since the purpose of this guide is to provide you with a process map to follow to work on the projects. This is an example here, you don't need to work on this, just follow a similar process to make a basic submission for any of the projects. Please note however that basic the submission that you make using the process mentioned here will not guarantee a score which is higher than the required threshold in individual projects. At the end of the guide, we will also discuss how to improve on your submissions.

```
In [1]: # List of commonly shortcuts that can be used in Jupyter notebook  
# Shift + Enter: Run the cell  
# Esc + A : Insert a cell above the current cell  
# Esc + B : Insert a cell below the current cell  
# Ctrl + / : Comment the selected code  
# Shift + Ctrl + - : Break the cell at the cursor position
```

Classification Problem - whether a person lies in a revenue.grid or not (1 or 0)

We have been provided with train and test data separately, where test data doesn't have response values. We need to use train data to build our model and then use that model to make prediction on test data and upload the csv file to LMS. In this we have been given data bd_train and bd_test and we are required to predict revenue.grid, whether it takes value 1 or not.

```
In [2]: import numpy as np  
import pandas as pd
```

```
In [3]: # Read the data  
bd_train = pd.read_csv(r"C:\Users\anjali\Dropbox\0.0 Data\bd_train.csv")
```

```
In [4]: bd_test = pd.read_csv(r"C:\Users\anjali\Dropbox\0.0 Data\bd_test.csv")
```

```
In [5]: bd_train.describe()
```

```
Out[5]:
```

	REF_NO	year_last_moved	Average.Credit.Card.Transaction	Balance.Transfer	Term.De
count	8124.000000	8124.000000	8124.000000	8124.000000	8124.00
mean	5750.928730	1967.511571	23.240174	46.369877	27.55
std	3316.648614	185.166494	50.714620	79.313730	53.98
min	1.000000	0.000000	0.000000	0.000000	0.00
25%	2900.750000	1978.000000	0.000000	0.000000	0.00
50%	5728.500000	1988.000000	0.000000	17.990000	0.00
75%	8621.500000	1994.000000	23.480000	65.452500	34.99
max	11518.000000	1999.000000	662.260000	2951.760000	784.82

```
In [6]: bd_train.head()
```

```
Out[6]:
```

	REF_NO	children	age_band	status	occupation	occupation_partner	home_status	fa
0	4888	Zero	55-60	Widowed	Retired	Unknown	Own Home	
1	8525	Zero	61-65	Partner	Retired	Retired	Own Home	
2	3411	3	31-35	Partner	Professional	Housewife	Own Home	
3	692	Zero	51-55	Partner	Secretarial/Admin	Other	Own Home	
4	10726	1	51-55	Partner	Retired	Retired	Own Home	

5 rows × 32 columns

Combine the train and test datasets

We will need same set of variables/features on both train and test datasets. It's easier to manage when we combine the train and test datasets in the beginning itself and then separate them once we are done with data preparation. Before combining however, we'll need some placeholder column which can be used to differentiate between the observations coming from train and test datasets. Also, we will need to add a column for the response variable to the test dataset in order to have the same columns in both train and test datasets. We'll fill test datasets response column with nan's.

```
In [7]: bd_train.shape
```

```
Out[7]: (8124, 32)
```

```
In [8]: bd_test.shape
```

```
Out[8]: (2031, 31)
```

```
In [9]: bd_train.columns
```

```
Out[9]: Index(['REF_NO', 'children', 'age_band', 'status', 'occupation',
              'occupation_partner', 'home_status', 'family_income', 'self_employed',
              'self_employed_partner', 'year_last_moved', 'TVarea', 'post_code',
              'post_area', 'Average.Credit.Card.Transaction', 'Balance.Transfer',
              'Term.Deposit', 'Life.Insurance', 'Medical.Insurance',
              'Average.A.C.Balance', 'Personal.Loan', 'Investment.in.Mutual.Fund',
              'Investment.Tax.Saving.Bond', 'Home.Loan', 'Online.Purchase.Amount',
              'Revenue.Grid', 'gender', 'region', 'Investment.in.Commudity',
              'Investment.in.Equity', 'Investment.in.Derivative',
              'Portfolio.Balance'],
             dtype='object')
```

```
In [10]: bd_test.columns # Revenue.Grid outcome variable not present in bd_test
```

```
Out[10]: Index(['REF_NO', 'children', 'age_band', 'status', 'occupation',
               'occupation_partner', 'home_status', 'family_income', 'self_employed',
               'self_employed_partner', 'year_last_moved', 'TVarea', 'post_code',
               'post_area', 'Average.Credit.Card.Transaction', 'Balance.Transfer',
               'Term.Deposit', 'Life.Insurance', 'Medical.Insurance',
               'Average.A.C.Balance', 'Personal.Loan', 'Investment.in.Mutual.Fund',
               'Investment.Tax.Saving.Bond', 'Home.Loan', 'Online.Purchase.Amount',
               'gender', 'region', 'Investment.in.Commudity', 'Investment.in.Equity',
               'Investment.in.Derivative', 'Portfolio.Balance'],
              dtype='object')
```

bd_test does not have the outcome variable 'Revenue.Grid' - in order to combine the two we will add the response variable to ld_test

```
In [11]: bd_test['Revenue.Grid']=np.nan
```

```
In [12]: # combine the two datasets to pre-process the data together since the data on
          which the model is trained should
          # undergo the same preprocessing as the data on which the predictions are made
```

```
In [13]: bd_train['data'] = 'train'
          bd_test['data'] = 'test'
          bd_test=bd_test[bd_train.columns] # the columns in the two data frames should
          be in the same order to enable concatenation
          bd_all=pd.concat([bd_train,bd_test],axis=0)
```

```
In [14]: bd_all.head()
```

Out[14]:

	REF_NO	children	age_band	status	occupation	occupation_partner	home_status	fa
0	4888	Zero	55-60	Widowed	Retired	Unknown	Own Home	
1	8525	Zero	61-65	Partner	Retired	Retired	Own Home	
2	3411	3	31-35	Partner	Professional	Housewife	Own Home	
3	692	Zero	51-55	Partner	Secretarial/Admin	Other	Own Home	
4	10726	1	51-55	Partner	Retired	Retired	Own Home	

5 rows × 33 columns

```
In [15]: bd_all.shape
```

Out[15]: (10155, 33)

In [16]: `bd_all.dtypes`

```
Out[16]: REF_NO          int64
children        object
age_band        object
status          object
occupation      object
occupation_partner object
home_status     object
family_income   object
self_employed   object
self_employed_partner object
year_last_moved int64
TVarea          object
post_code       object
post_area       object
Average.Credit.Card.Transaction float64
Balance.Transfer float64
Term.Deposit     float64
Life.Insurance   float64
Medical.Insurance float64
Average.A.C.Balance float64
Personal.Loan    float64
Investment.in.Mutual.Fund float64
Investment.Tax.Saving.Bond float64
Home.Loan        float64
Online.Purchase.Amount float64
Revenue.Grid     float64
gender           object
region           object
Investment.in.Commudity float64
Investment.in.Equity float64
Investment.in.Derivative float64
Portfolio.Balance float64
data             object
dtype: object
```

In [17]: `bd_all.describe()`

```
Out[17]:
```

	REF_NO	year_last_moved	Average.Credit.Card.Transaction	Balance.Transfer	Term.De
count	10155.000000	10155.000000	10155.000000	10155.000000	10155.00
mean	5770.830822	1968.376366	23.441757	46.417760	27.57
std	3324.837813	180.202242	50.872127	78.477609	53.95
min	1.000000	0.000000	0.000000	0.000000	0.00
25%	2903.500000	1978.000000	0.000000	0.000000	0.00
50%	5770.000000	1988.000000	0.000000	17.960000	0.00
75%	8665.500000	1994.000000	23.980000	65.385000	34.95
max	11518.000000	1999.000000	662.260000	2951.760000	784.82

Handling each variable

```
In [18]: list(zip(bd_all.columns, bd_all.dtypes, bd_all.nunique())) # gives the type and
the number of unique values present for each column
```

```
Out[18]: [('REF_NO', dtype('int64'), 10155),
('children', dtype('O'), 5),
('age_band', dtype('O'), 13),
('status', dtype('O'), 5),
('occupation', dtype('O'), 9),
('occupation_partner', dtype('O'), 9),
('home_status', dtype('O'), 5),
('family_income', dtype('O'), 13),
('self_employed', dtype('O'), 2),
('self_employed_partner', dtype('O'), 2),
('year_last_moved', dtype('int64'), 95),
('TVarea', dtype('O'), 14),
('post_code', dtype('O'), 10040),
('post_area', dtype('O'), 2039),
('Average.Credit.Card.Transaction', dtype('float64'), 1411),
('Balance.Transfer', dtype('float64'), 2183),
('Term.Deposit', dtype('float64'), 1419),
('Life.Insurance', dtype('float64'), 3111),
('Medical.Insurance', dtype('float64'), 1589),
('Average.A.C.Balance', dtype('float64'), 2223),
('Personal.Loan', dtype('float64'), 1760),
('Investment.in.Mutual.Fund', dtype('float64'), 2470),
('Investment.Tax.Saving.Bond', dtype('float64'), 832),
('Home.Loan', dtype('float64'), 884),
('Online.Purchase.Amount', dtype('float64'), 1319),
('Revenue.Grid', dtype('float64'), 2),
('gender', dtype('O'), 3),
('region', dtype('O'), 13),
('Investment.in.Commudity', dtype('float64'), 3558),
('Investment.in.Equity', dtype('float64'), 3250),
('Investment.in.Derivative', dtype('float64'), 3796),
('Portfolio.Balance', dtype('float64'), 8317),
('data', dtype('O'), 2)]
```

```
In [19]: # REF_NO, post_code , post_area : drop
# children : Zero : 0 , 4+ : 4 and then convert to numeric
# age_band : dummies
# status , occupation , occupation_partner , home_status, family_income : dummies
# self_employed, ` : dummies
# TVArea , Region , gender : dummies
# Revenue Grid : 1, 2 : 1, 0
```

```
In [20]: bd_all['children'].value_counts()
```

```
Out[20]: Zero      6208  
         1       1848  
         2       1607  
         3        473  
         4+        19  
         Name: children, dtype: int64
```

```
In [21]: # children : Zero : 0 , 4+ : 4 and then convert to numeric  
bd_all['children']=np.where(bd_all['children']=='Zero',0,bd_all['children']) #  
         replace 'Zero' with 0  
bd_all['children']=np.where(bd_all['children'][:1]=='4',4,bd_all['children'])  
         # replace '4' and '4+' with 4  
bd_all['children']=pd.to_numeric(bd_all['children'],errors='coerce')
```

If we look at variables `post_code` and `post_area`, you will notice that they have more than 2000 unique values, having none of the individual values having frequency higher than a good number , we'll drop those variables.

```
In [22]: print(bd_all['post_area'].unique()) # gives the unique values in the column  
print(bd_all['post_area'].nunique()) # gives the number of unique values  
bd_all['post_area'].value_counts() # gives counts of each unique value in the  
column
```



```
[ 'L15' 'CW11' 'PR8' ... 'P039' 'WA44' 'WV4' ]  
2039
```

```
Out[22]: PR5      35
          PR4      33
          WA7      28
          TQ12     28
          M12      26
          WA4      26
          LE7      25
          BH23     25
          LE9      24
          M33      24
          SK8      22
          PR8      22
          M13      21
          WA5      21
          M14      21
          TS5      21
          HD7      21
          WA3      20
          NG9      20
          PR2      20
          PR7      20
          PR6      20
          L6       20
          DY8      20
          L12      20
          L13      20
          CF62     20
          L8       19
          L15      19
          L40      19
          ..
          PL29     1
          HP8      1
          NW8      1
          IP2      1
          EX9      1
          DL4      1
          TW5      1
          BS3      1
          OL5      1
          BD7      1
          BH2      1
          IP21     1
          CA9      1
          BD14     1
          WV4      1
          AB13     1
          NE17     1
          W10      1
          P016     1
          OX7      1
          S024     1
          N12      1
          TN21     1
          CA16     1
          LL57     1
          PL27     1
```

PL9	1
SW16	1
MK19	1
TN17	1

Name: post_area, Length: 2039, dtype: int64

```
In [23]: print(bd_all['post_code'].unique())  
         print(bd_all['post_code'].nunique())  
         bd_all['post_code'].value_counts()
```

```
['L15 6LL' 'CW11 1RA' 'PR8 6SQ' ... 'CF43 4SU' 'IV30 6AX' 'SM4 5RF']  
10040
```

```

Out[23]: SA2 0FN      2
         MK7 7HH      2
         Y08 9JZ      2
         BD20 8BY     2
         L31 6NN      2
         CH3 9HA      2
         S41 0SU      2
         NP13 1QD     2
         CF48 1EW     2
         EN4 8QQ      2
         SN14 6PS     2
         M14 7NQ      2
         KY11 9GF     2
         CW1 5TY      2
         TQ14 8RZ     2
         WS9 9RD      2
         DT9 6LA      2
         LE65 1RN     2
         L9 7BN       2
         IG9 5DZ      2
         CF3 4DF      2
         CM8 2QR      2
         CM6 2JA      2
         GL52 6SX     2
         CV32 6HE     2
         AB23 8QW     2
         HU9 4BT      2
         BH1 1QG      2
         TQ12 6YA     2
         LU7 7UQ      2
         ..
         S72 8HJ      1
         IV15 9TU     1
         CH2 1QH      1
         WA4 5HP      1
         M4 8LZ       1
         CF62 6LZ     1
         B97 5HF      1
         M46 0AY      1
         EX31 2NY     1
         L53 3NE      1
         WA4 2LF      1
         TQ9 7ES      1
         LU3 3HF      1
         L1 4DN       1
         SA11 3TQ     1
         DT11 0HW     1
         M14 5ES      1
         RG26 3YN     1
         DE56 0HX     1
         DT9 6PT      1
         CF10 1DL     1
         SP10 4LP     1
         TS12 1PE     1
         Y032 9QA     1
         NE33 5SZ     1
         CF45 4LD     1

```

```
G62 6QL      1
RG5 4TJ      1
TN36 4BJ     1
BL0 9LH      1
Name: post_code, Length: 10040, dtype: int64
```

```
In [24]: # Dropping the variables Post.Code and Post.Area
bd_all.drop(['post_code', 'post_area'], axis=1, inplace=True)
```

```
In [25]: # Dropping Ref_no column since it has no use in building the model
bd_all.drop('REF_NO', axis=1, inplace=True)
```

```
In [26]: # We will create dummies for the rest of the categorical columns using the built-in function from pandas
```

```
In [27]: cat_vars=bd_all.select_dtypes(['object']).columns # cat_vars variable contains all the categorical variables
cat_vars # we will create dummies for all of them using the built-in pandas library
```

```
Out[27]: Index(['age_band', 'status', 'occupation', 'occupation_partner', 'home_status',
               'family_income', 'self_employed', 'self_employed_partner', 'TVarea',
               'gender', 'region', 'data'],
              dtype='object')
```

```
In [28]: for col in cat_vars[:-1]: # we do not consider the 'data' column since that is only a placeholder for train and test sets
        dummy=pd.get_dummies(bd_all[col],drop_first=True,prefix=col)
        bd_all=pd.concat([bd_all,dummy],axis=1)
        del bd_all[col]
        print(col) # gives the columns for which the dummies are created and which have been deleted
del dummy
# we use a for loop to create dummies for each categorical column/ feature/ variable
```

```
age_band
status
occupation
occupation_partner
home_status
family_income
self_employed
self_employed_partner
TVarea
gender
region
```

You can also use the method to create dummies as described in the Project Process Guide for Regression - Python. That functions enables you to ignore categories with low frequencies

```
In [29]: # we could also have created dummies separately for each variable as follows:
# dummy1=pd.get_dummies(bd_all['age_band'],drop_first=True,prefix='age_band')
# del bd_all['age_band']
# dummy2=pd.get_dummies(bd_all['status'],drop_first=True,prefix='status')
# del bd_all['status']
# dummy3=pd.get_dummies(bd_all['occupation'],drop_first=True,prefix='occupatio
n')
# del bd_all['occupation']
# dummy4=pd.get_dummies(bd_all['occupation_partner'],drop_first=True,prefix='o
ccupation_partner')
# del bd_all['occupation_partner']
# dummy5=pd.get_dummies(bd_all['home_status'],drop_first=True,prefix='home_sta
tus')
# del bd_all['home_status']
# dummy6=pd.get_dummies(bd_all['family_income'],drop_first=True,prefix='family
_income')
# del bd_all['family_income']
# dummy7=pd.get_dummies(bd_all['self_employed'],drop_first=True,prefix='self_e
mployed')
# del bd_all['self_employed']
# dummy8=pd.get_dummies(bd_all['self_employed_partner'],drop_first=True,prefix
='self_employed_partner')
# del bd_all['self_employed_partner']
# dummy9=pd.get_dummies(bd_all['TVarea'],drop_first=True,prefix='TVarea')
# del bd_all['TVarea']
# dummy10=pd.get_dummies(bd_all['gender'],drop_first=True,prefix='gender')
# del bd_all['gender']
# dummy11=pd.get_dummies(bd_all['region'],drop_first=True,prefix='region')
# del bd_all['region']
# test1 = pd.concat([bd_all,dummy1,dummy2,dummy3,dummy4,dummy5,dummy6,dummy7,d
ummy8,dummy9,dummy10,dummy11],axis=1)
# test1.columns
# len(test1.columns)
```

```
In [30]: bd_all.shape
```

```
Out[30]: (10155, 96)
```



```
In [31]: bd_all.dtypes # check if all the columns are numeric
```

```

Out[31]: children                float64
        year_last_moved          int64
        Average.Credit.Card.Transaction float64
        Balance.Transfer         float64
        Term.Deposit             float64
        Life.Insurance           float64
        Medical.Insurance        float64
        Average.A.C.Balance      float64
        Personal.Loan           float64
        Investment.in.Mutual.Fund float64
        Investment.Tax.Saving.Bond float64
        Home.Loan               float64
        Online.Purchase.Amount   float64
        Revenue.Grid            float64
        Investment.in.Commudity   float64
        Investment.in.Equity      float64
        Investment.in.Derivative  float64
        Portfolio.Balance        float64
        data                    object
        age_band_22-25           uint8
        age_band_26-30           uint8
        age_band_31-35           uint8
        age_band_36-40           uint8
        age_band_41-45           uint8
        age_band_45-50           uint8
        age_band_51-55           uint8
        age_band_55-60           uint8
        age_band_61-65           uint8
        age_band_65-70           uint8
        age_band_71+            uint8
        ...
        family_income_Unknown    uint8
        self_employed_Yes        uint8
        self_employed_partner_Yes uint8
        TVarea_Border            uint8
        TVarea_Carlton           uint8
        TVarea_Central           uint8
        TVarea_Grampian          uint8
        TVarea_Granada           uint8
        TVarea_HTV               uint8
        TVarea_Meridian          uint8
        TVarea_Scottish TV       uint8
        TVarea_TV South West     uint8
        TVarea_Tyne Tees         uint8
        TVarea_Ulster            uint8
        TVarea_Unknown           uint8
        TVarea_Yorkshire         uint8
        gender_Male              uint8
        gender_Unknown           uint8
        region_East Anglia        uint8
        region_East Midlands     uint8
        region_Isle of Man       uint8
        region_North             uint8
        region_North West        uint8
        region_Northern Ireland  uint8
        region_Scotland          uint8
        region_South East        uint8

```

```
region_South West      uint8
region_Unknown         uint8
region_Wales           uint8
region_West Midlands   uint8
Length: 96, dtype: object
```

```
In [32]: # Analyzing the outcome variable
bd_all['Revenue.Grid'].value_counts()
```

```
Out[32]: 2.0    7256
         1.0     868
         Name: Revenue.Grid, dtype: int64
```

```
In [33]: # Change the response variable value to '1' and '0' from '2' and '1'
bd_all['Revenue.Grid']=(bd_all['Revenue.Grid']==1).astype(int)
```

Check for missing values

In [34]: `bd_all.isnull().sum()`

```

Out[34]: children          19
         year_last_moved    0
         Average.Credit.Card.Transaction  0
         Balance.Transfer    0
         Term.Deposit        0
         Life.Insurance      0
         Medical.Insurance   0
         Average.A.C.Balance  0
         Personal.Loan       0
         Investment.in.Mutual.Fund  0
         Investment.Tax.Saving.Bond  0
         Home.Loan           0
         Online.Purchase.Amount  0
         Revenue.Grid        0
         Investment.in.Commudity  0
         Investment.in.Equity  0
         Investment.in.Derivative  0
         Portfolio.Balance    0
         data                0
         age_band_22-25      0
         age_band_26-30      0
         age_band_31-35      0
         age_band_36-40      0
         age_band_41-45      0
         age_band_45-50      0
         age_band_51-55      0
         age_band_55-60      0
         age_band_61-65      0
         age_band_65-70      0
         age_band_71+        0
         ..
         family_income_Unknown  0
         self_employed_Yes      0
         self_employed_partner_Yes  0
         TVarea_Border          0
         TVarea_Carlton         0
         TVarea_Central         0
         TVarea_Grampian        0
         TVarea_Granada         0
         TVarea_HTV             0
         TVarea_Meridian        0
         TVarea_Scottish TV     0
         TVarea_TV South West   0
         TVarea_Tyne Tees       0
         TVarea_Ulster          0
         TVarea_Unknown         0
         TVarea_Yorkshire       0
         gender_Male            0
         gender_Unknown         0
         region_East Anglia     0
         region_East Midlands   0
         region_Isle of Man     0
         region_North           0
         region_North West     0
         region_Northern Ireland  0
         region_Scotland        0
         region_South East      0

```

```
region_South West      0
region_Unknown          0
region_Wales            0
region_West Midlands    0
Length: 96, dtype: int64
```

```
In [35]: # Only the feature 'children' has missing values. We ignore the missing values  
         in the response variable 'Revenue.Grid'  
# Handling missing values in 'children' taking the mean only from the train da  
taset  
bd_all.loc[bd_all['children'].isnull(), 'children'] = bd_all.loc[bd_all['data'] ==  
    'train', 'children'].mean()
```

```
In [36]: # Check again if there are any missing values  
bd_all.isnull().sum()
```

```

Out[36]: children                                0
         year_last_moved                        0
         Average.Credit.Card.Transaction        0
         Balance.Transfer                       0
         Term.Deposit                           0
         Life.Insurance                         0
         Medical.Insurance                      0
         Average.A.C.Balance                    0
         Personal.Loan                         0
         Investment.in.Mutual.Fund              0
         Investment.Tax.Saving.Bond             0
         Home.Loan                             0
         Online.Purchase.Amount                 0
         Revenue.Grid                           0
         Investment.in.Commudity                0
         Investment.in.Equity                   0
         Investment.in.Derivative               0
         Portfolio.Balance                     0
         data                                  0
         age_band_22-25                        0
         age_band_26-30                        0
         age_band_31-35                        0
         age_band_36-40                        0
         age_band_41-45                        0
         age_band_45-50                        0
         age_band_51-55                        0
         age_band_55-60                        0
         age_band_61-65                        0
         age_band_65-70                        0
         age_band_71+                          0
         ..
         family_income_Unknown                 0
         self_employed_Yes                     0
         self_employed_partner_Yes              0
         TVarea_Border                         0
         TVarea_Carlton                       0
         TVarea_Central                       0
         TVarea_Grampian                      0
         TVarea_Granada                       0
         TVarea_HTV                           0
         TVarea_Meridian                      0
         TVarea_Scottish TV                    0
         TVarea_TV South West                  0
         TVarea_Tyne Tees                     0
         TVarea_Ulster                        0
         TVarea_Unknown                       0
         TVarea_Yorkshire                     0
         gender_Male                           0
         gender_Unknown                       0
         region_East Anglia                    0
         region_East Midlands                  0
         region_Isle of Man                   0
         region_North                         0
         region_North West                    0
         region_Northern Ireland              0
         region_Scotland                      0
         region_South East                    0

```



```

region_South West      0
region_Unknown         0
region_Wales           0
region_West Midlands   0
Length: 96, dtype: int64

```

In [37]: *# Data preprocessing is complete - the data is in the expected format*

Let's separate our two data sets and remove the unnecessary columns that we added while combining them.

In [38]: `bd_train=bd_all[bd_all['data']=='train']` *# Select only those rows where the data column has the value 'train'*
`del bd_train['data']` *# Remove the data column from bd_train data frame*

In [39]: `bd_test=bd_all[bd_all['data']=='test']` *# Select only those rows where the data column has the value 'test'*
`bd_test.drop(['Revenue.Grid', 'data'], axis=1, inplace=True)` *# Remove the data and Interest.Rate column from bd_test data frame*

C:\Users\anjali\Anaconda3\lib\site-packages\pandas\core\frame.py:3697: Setting WithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
errors=errors)

In [40]: `del bd_all` *# After data pre-processing, we do not need the bd_all data frame*

Model Implementation

Let's build a model on training data. We are not going to build any complex model, it's up to you to use other algorithms that you learn during the course.

In [41]: *# In the project process guide for regression, we broke the train data further into train and validation datasets to enable us to assess the models performance*
Instead, here we will use cross validation to get the best parameters for the model and assess model performance

Logistic Regression

In [42]: `from sklearn.linear_model import LogisticRegression`

In [43]: `params={'class_weight':['balanced',None],`
`'penalty':['l1','l2'],`
`'C':np.linspace(0.01,1000,10)}`

```
In [44]: x_train=bd_train.drop('Revenue.Grid',axis=1)
y_train=bd_train['Revenue.Grid']
# y_train.dtype
```

```
In [45]: model=LogisticRegression(fit_intercept=True)
```

```
In [46]: from sklearn.model_selection import GridSearchCV
```

```
In [47]: grid_search=GridSearchCV(model,param_grid=params,cv=5,scoring="roc_auc")
```

```
In [48]: grid_search.fit(x_train,y_train) # grid search fits a Logistic Regression model for every combination of parameters
# we can then select the parameter combination that gives the best performance
# Randomized search on the other hand would take a subset of parameter combinations and fit the model. Algorithms in which many parameters need to be tuned, randomized search makes more sense
```

```
Out[48]: GridSearchCV(cv=5, error_score='raise',
    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False),
    fit_params=None, iid=True, n_jobs=1,
    param_grid={'class_weight': ['balanced', None], 'penalty': ['l1', 'l2'], 'C': array([1.0000e-02, 1.1112e+02, 2.2223e+02, 3.3334e+02, 4.4445e+02, 5.5556e+02, 6.6667e+02, 7.7778e+02, 8.8889e+02, 1.0000e+03])},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=0)
```

```
In [49]: grid_search.best_estimator_ # returns the parameters giving the best performance
# grid search with cv chose C as 0.01, class_weight as 'balanced' and penalty as 'l1'
```

```
Out[49]: LogisticRegression(C=0.01, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l1', random_state=None,
    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

```
In [50]: logr=grid_search.best_estimator_
```

Using the report function given below you can see the cv performance of top few models that will be the tentative performance

```
In [51]: def report(results, n_top=3):
        for i in range(1, n_top + 1):
            candidates = np.flatnonzero(results['rank_test_score'] == i)
            for candidate in candidates:
                print("Model with rank: {0}".format(i))
                print("Mean validation score: {0:.3f} (std: {1:.3f})".format(
                    results['mean_test_score'][candidate],
                    results['std_test_score'][candidate]))
                print("Parameters: {0}".format(results['params'][candidate]))
            print("")
```

```
In [52]: report(grid_search.cv_results_,5)
```

```
Model with rank: 1
Mean validation score: 0.955 (std: 0.005)
Parameters: {'C': 0.01, 'class_weight': 'balanced', 'penalty': 'l1'}

Model with rank: 2
Mean validation score: 0.955 (std: 0.005)
Parameters: {'C': 0.01, 'class_weight': 'balanced', 'penalty': 'l2'}

Model with rank: 3
Mean validation score: 0.951 (std: 0.005)
Parameters: {'C': 888.89, 'class_weight': 'balanced', 'penalty': 'l2'}

Model with rank: 4
Mean validation score: 0.951 (std: 0.005)
Parameters: {'C': 555.56, 'class_weight': 'balanced', 'penalty': 'l2'}

Model with rank: 5
Mean validation score: 0.951 (std: 0.005)
Parameters: {'C': 444.45, 'class_weight': 'balanced', 'penalty': 'l2'}
```

```
In [53]: logr.fit(x_train,y_train) # we fit the model with the best parameters
```

```
Out[53]: LogisticRegression(C=0.01, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='ovr', n_jobs=1, penalty='l1', random_state=None,
                             solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

```
In [54]: train_score=logr.predict_proba(x_train)[:,:1] # we can get the probabilities using predict_proba
```

```
In [55]: train_score
```

```
Out[55]: array([0.99420408, 0.01137188, 0.06167701, ..., 0.80850512, 0.17254272,
                0.03292699])
```

In case you have to submit probabilities, please use the following code:

```
In [56]: test_score=logr.predict_proba(bd_test)[: ,1]
         test_score
```

```
Out[56]: array([0.25083295, 0.38726091, 0.00597118, ..., 0.00525166, 0.04848874,
                0.11929218])
```

```
In [57]: pd.DataFrame(test_score).to_csv("mysubmission.csv", index=False)
```

Upload this csv file to LMS under that project submission tab

However, we may want hard classes instead of probabilities

```
In [58]: cutoffs=np.linspace(0.01,0.99,99)
         cutoffs
```

```
Out[58]: array([0.01, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08, 0.09, 0.1 , 0.11,
                0.12, 0.13, 0.14, 0.15, 0.16, 0.17, 0.18, 0.19, 0.2 , 0.21, 0.22,
                0.23, 0.24, 0.25, 0.26, 0.27, 0.28, 0.29, 0.3 , 0.31, 0.32, 0.33,
                0.34, 0.35, 0.36, 0.37, 0.38, 0.39, 0.4 , 0.41, 0.42, 0.43, 0.44,
                0.45, 0.46, 0.47, 0.48, 0.49, 0.5 , 0.51, 0.52, 0.53, 0.54, 0.55,
                0.56, 0.57, 0.58, 0.59, 0.6 , 0.61, 0.62, 0.63, 0.64, 0.65, 0.66,
                0.67, 0.68, 0.69, 0.7 , 0.71, 0.72, 0.73, 0.74, 0.75, 0.76, 0.77,
                0.78, 0.79, 0.8 , 0.81, 0.82, 0.83, 0.84, 0.85, 0.86, 0.87, 0.88,
                0.89, 0.9 , 0.91, 0.92, 0.93, 0.94, 0.95, 0.96, 0.97, 0.98, 0.99])
```

```
In [59]: train_score=logr.predict_proba(x_train)[: ,1] # the predicted response variable
         values
         real=y_train # the actual response variable values
         print(logr.classes_) # In order to find the probability of which column is fo
         r outcome 1 and which for outcome 0

         [0 1]
```

```
In [60]: from sklearn.metrics import fbeta_score
```

```
In [61]: KS_all=[]

for cutoff in cutoffs:

    predicted=(train_score>cutoff).astype(int)

    TP=((predicted==1) & (real==1)).sum()
    TN=((predicted==0) & (real==0)).sum()
    FP=((predicted==1) & (real==0)).sum()
    FN=((predicted==0) & (real==1)).sum()

    P=TP+FN
    N=TN+FP

    KS=(TP/P)-(FP/N)

    KS_all.append(KS)

# try out what cutoffs you get when you use F_beta scores with different v
alues of betas [0.5 , 5]
# beta < 1 : you will get cutoff , which is high ( favours precision)
# beta > 1 : you will get cutoff , which is low (favours precision )
```

```
In [62]: list(zip(cutoffs,KS_all))
```

```

Out[62]: [(0.01, 0.10276192847235277),
(0.02, 0.1574644724340638),
(0.03, 0.2101988375106062),
(0.04, 0.25980215324739997),
(0.05, 0.30922891082669857),
(0.060000000000000005, 0.3677515890234174),
(0.06999999999999999, 0.43371638408893454),
(0.08, 0.5118090733110117),
(0.09, 0.5500231176868087),
(0.09999999999999999, 0.583088395937384),
(0.11, 0.6104148989680873),
(0.12, 0.6311261870043035),
(0.13, 0.6496819412759948),
(0.14, 0.6632267464015161),
(0.15000000000000002, 0.6780614422388083),
(0.16, 0.6912423343274785),
(0.17, 0.7019425207932162),
(0.18000000000000002, 0.7109393656100275),
(0.19, 0.7192471255315798),
(0.2, 0.7271414345159766),
(0.21000000000000002, 0.7350852814006777),
(0.22, 0.7435803962015862),
(0.23, 0.7490435374633546),
(0.24000000000000002, 0.7535807010502035),
(0.25, 0.7599202820865871),
(0.26, 0.7634044477413258),
(0.27, 0.7672137852544725),
(0.28, 0.7708853057885672),
(0.29000000000000004, 0.7717014109410169),
(0.3, 0.7749982217163993),
(0.31, 0.7802352669203685),
(0.32, 0.7843202383916187),
(0.33, 0.7860127833186837),
(0.34, 0.7896843038527784),
(0.35000000000000003, 0.7917902997169989),
(0.36000000000000004, 0.792330770911345),
(0.37, 0.7919948023310758),
(0.38, 0.796404945660734),
(0.39, 0.7997125531579776),
(0.4, 0.798963133640553),
(0.41000000000000003, 0.8015321183422331),
(0.42000000000000004, 0.8051153597975805),
(0.43, 0.8059810028503346),
(0.44, 0.8031643286471326),
(0.45, 0.8060089473069165),
(0.46, 0.8083022980504931),
(0.47000000000000003, 0.8079275882917807),
(0.48000000000000004, 0.8093940371610464),
(0.49, 0.8066161041362876),
(0.5, 0.8064674904353746),
(0.51, 0.8096372809535665),
(0.52, 0.8053433611592378),
(0.53, 0.805244285358629),
(0.54, 0.8052830265370721),
(0.55, 0.8063360244691823),
(0.56, 0.7978967985814377),
(0.57000000000000001, 0.7983489906970364),

```

```
(0.5800000000000001, 0.7946558767192191),
(0.59, 0.7906871287832984),
(0.6, 0.7911393208988969),
(0.61, 0.7856050482931018),
(0.62, 0.7833396420061071),
(0.63, 0.7807986017610088),
(0.64, 0.7814381487559636),
(0.65, 0.7790844633902214),
(0.66, 0.7739636417215817),
(0.67, 0.7665386725875043),
(0.68, 0.7641354493214578),
(0.6900000000000001, 0.7590146276528181),
(0.7000000000000001, 0.7524660982933559),
(0.7100000000000001, 0.7469318256875607),
(0.72, 0.7422244549560764),
(0.73, 0.7312054476447903),
(0.74, 0.7268232487717141),
(0.75, 0.7211016212865629),
(0.76, 0.7129875672572261),
(0.77, 0.7093439911797134),
(0.78, 0.7037601806736138),
(0.79, 0.6917269166086607),
(0.8, 0.6838389586371235),
(0.81, 0.6789937709265874),
(0.8200000000000001, 0.6741981211163556),
(0.8300000000000001, 0.6637799196215813),
(0.8400000000000001, 0.6513332046194727),
(0.85, 0.6420670768574172),
(0.86, 0.6304968016299239),
(0.87, 0.61763663569066),
(0.88, 0.6132544368175836),
(0.89, 0.6082714321279957),
(0.9, 0.5906651542788044),
(0.91, 0.5752252069159989),
(0.92, 0.5502434978330344),
(0.93, 0.5326372199838431),
(0.9400000000000001, 0.519451882186171),
(0.9500000000000001, 0.49672478266833997),
(0.9600000000000001, 0.4740472210508132),
(0.97, 0.44182789771312725),
(0.98, 0.4103471971710048),
(0.99, 0.3536586914881185)]
```

```
In [63]: mycutoff=cutoffs[KS_all==max(KS_all)][0]
mycutoff # gives the cutoff value where KS is maximum
```

```
Out[63]: 0.51
```

```
In [64]: test_score=logr.predict_proba(bd_test)[:,1]
test_score
```

```
Out[64]: array([0.25083295, 0.38726091, 0.00597118, ..., 0.00525166, 0.04848874,
0.11929218])
```


If you had to submit hardclasses, you can apply the cutoff obtained above and then submit

```
In [65]: test_classes=(test_score>mycutoff).astype(int)

In [66]: pd.DataFrame(test_classes).to_csv("mysubmission.csv",index=False)
         # Please give the file a proper name before submission
```

This csv file is what you need to upload on LMS.

Passing the project

You'll clear a project if you clear the threshold. Every project has a different threshold, clearly mentioned in the project problem statement page. Do go through the problem statement page and find that out before assuming that you cleared a project just by making a submission. You are free to make as many submissions as you want; maximum score among all your submissions will be considered as your latest score.

Improving your score

Try a non linear algorithm such as dtrees, random forests, extratrees or gbm; they almost always give better performance than a simple linear model. Tune your parameters for above mentioned algorithms. Try stacking. Please remember to do one thing at a time.

How do I assess my model performance before submission

You can break your given train data into two parts, build model on one check its performance on the other. That will give you an idea about how your model might perform after submission. You can not assess performance on given test data because response has been removed from it. Make sure that the submission that you make is based on the model built on entire training data.

Do i need to do this all by myself?

Projects in the course are an integrated learning process. If you get stuck somewhere, feel free to ask for help on QA forum. However, don't expect outright solutions, such requests will be ignored and deleted from QA forum. Hope this helps; feel free to suggest if we need to include anything else in here.